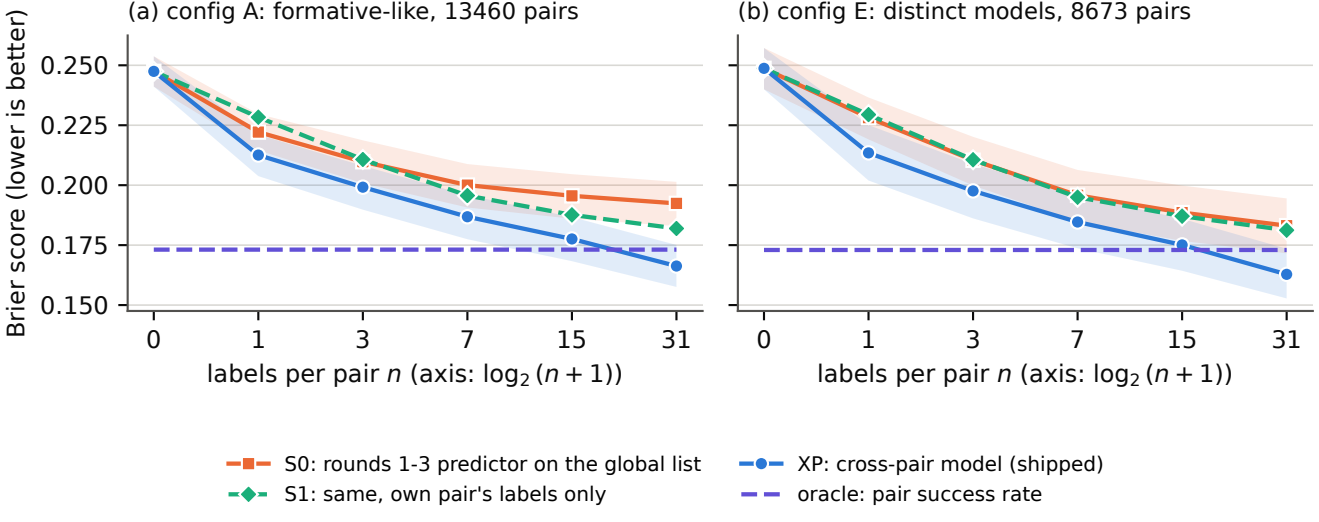
(a) config A: formative-like, 13460 pairs
(b) config E: distinct models, 8673 pairs
Brier score (lower is better)
0.250
0.225
0.200
0.175
0.150
0
1
3
7
15
31
labels per pair n (axis: $\log_2(n+1)$)
S0: rounds 1-3 predictor on the global list
S1: same, own pair's labels only
XP: cross-pair model (shipped)
oracle: pair success rate